



# Basic String Manipulation and Regular Expression

ITCS 209

Assistant Prof. Dr. Suppawong Tuarob

Faculty of Information and Communication Technology





# The String Class

- ◆ Strings are "**immutable**" objects: once instantiated, a String object is constant and not changeable.
- ◆ because String objects are immutable, they can be shared fearlessly: no one can change your object.
- ◆ Java optimizes memory by maintaining a pool of shared Strings: the constants "A", "A", "A" will have three references to the same String object.
- ◆ Note: an empty string object is not null
  - ◆ literal: "" is a String object with a length() of 0





# Strings are special in Java

- ▶ Strings are used so often in programming, Java makes special allowances for coding them
- ▶ construct a String without new String()  
`String s1 = new String("some text");`  
`String s2 = "more text";`
- ▶ the only overloaded operators in Java are for Strings
- ▶ `s1 += s2; // s1 = s1.concat(s2);`
- ▶ `s1 = s2 + "etc"; // s1 = s2.concat("etc");`





# Anything can be a String, just ask

- ▶ static method `String.valueOf()`
  - ▶ returns a `String`
  - ▶ can take almost anything as a parameter: all primitives, char array, any object.
- ▶ all objects inherit or override the `Object` class `toString()` method
- ▶ `System.out.println()` automatically calls `toString()` on any object in the parameter list
  - ▶ `System.out.println(myObject);` // is same as
  - ▶ `System.out.println(myObject.toString() );`





# The String Class

- ◆ comparison of two String objects:

- ◆ `thisString.equals(thatString)` *compares contents*
  - ◆ *this is what most of us mean most of the time*
- ◆ `thisString == thatString` *compares obj.ref.*
  - ◆ *this may seem like it works but is unreliable*





# Useful String Class Methods

- ◆ `length()` returns int of character count
- ◆ `trim()` returns String exclusive of lead/trail blanks
- ◆ `toUpperCase()`, `toLowerCase()` returns consistent case
- ◆ `valueOf()` returns String of any primitive
- ◆ `indexOf()` returns int locating a char or substring
- ◆ `charAt()` allows processing of string like `char[]`
- ◆ `substring()` returns a substring from this string
- ◆ `replace()` changes characters
- ◆ `replaceAll()` changes strings with regular expressions
- ◆ `split()` splits a string into an array of strings using `reg.exp.`





# StringBuilder

- ▶ StringBuilder is a mutable version of String.
- ▶ You can keep `append()`'ing without having to create a new string object.
- ▶ Save a lot of time and memory when dealing with massive text.
- ▶ For example, if you need to process and write 1M double values into a text file one by one.
  - ▶ It's faster to keep appending them to a StringBuilder object first, then write the whole thing into the file.
  - ▶ Disk I/O is an expensive operation.





# Regular Expressions in Java







# Regular Expressions

- ▶ A regular expression is a kind of pattern that can be applied to text (**Strings**, in Java)
- ▶ A regular expression either matches the text (or part of the text), or it fails to match
  - ▶ If a regular expression matches a part of the text, then you can easily find out which part
  - ▶ If a regular expression is complex, then you can easily find out which parts of the regular expression match which parts of the text
  - ▶ With this information, you can readily extract parts of the text, or do substitutions in the text
- ▶ Regular expressions are an extremely useful tool for manipulating text
  - ▶ Regular expressions are heavily used in the automatic generation of Web pages





## A first example

- ▶ The regular expression "[a-z]+" will match a sequence of one or more lowercase letters

[a-z] means any character from a through z, inclusive  
+ means "one or more"

- ▶ Suppose we apply this pattern to the String "Now is the time"

- ▶ There are *three ways* we can apply this pattern:

- ▶ To the *entire string*: it fails to match because the string contains characters other than lowercase letters
- ▶ To the *beginning of the string*: it fails to match because the string does not begin with a lowercase letter
- ▶ To *search the string*: it will succeed and match **ow**
  - ▶ If the pattern is applied a second time, it will find **is**
  - ▶ Further applications will find **is**, then **the**, then **time**
  - ▶ After **time**, another application will fail





# Doing it in Java, I

- First, you must *compile* the pattern

```
import java.util.regex.*;  
Pattern p = Pattern.compile("[a-z]+");
```

- Next, you must create a *matcher* for a specific piece of text by sending a message to your pattern

```
Matcher m = p.matcher("Now is the time");
```

- Points to notice:

- **Pattern** and **Matcher** are both in `java.util.regex`
- Neither **Pattern** nor **Matcher** has a public constructor; you create these by using methods in the **Pattern** class
- The matcher contains information about *both* the pattern to use *and* the text to which it will be applied





# Doing it in Java, II

■ Now that we have a matcher `m`,

- `m.matches()` returns `true` if the pattern matches the entire text string, and `false` otherwise
- `m.looksAt()` returns `true` if the pattern matches at the beginning of the text string, and `false` otherwise
- `m.find()` returns `true` if the pattern matches any part of the text string, and `false` otherwise
  - If called again, `m.find()` will start searching from where the last match was found
  - `m.find()` will return `true` for as many matches as there are in the string; after that, it will return `false`
  - When `m.find()` returns `false`, matcher `m` will be *reset* to the beginning of the text string (and may be used again)





# Finding what was matched

- ▶ After a successful match, `m.start()` will return the index of the first character matched
- ▶ After a successful match, `m.end()` will return the index of the last character matched, *plus one*
- ▶ If no match was attempted, or if the match was unsuccessful, `m.start()` and `m.end()` will throw an `IllegalStateException`
  - ▶ This is a `RuntimeException`, so you don't have to catch it
- ▶ It may seem strange that `m.end()` returns the index of the last character matched plus one, but this is just what most String methods require
  - ▶ For example, `"Now is the time".substring(m.start(), m.end())` will return exactly the matched substring





# A complete example

```
► import java.util.regex.*;
```

```
public class RegexTest {  
    public static void main(String args[]) {  
        String pattern = "[a-z]+";  
        String text = "Now is the time";  
        Pattern p = Pattern.compile(pattern);  
        Matcher m = p.matcher(text);  
        while (m.find()) {  
            System.out.print(text.substring(m.start(),  
                                           m.end()) + "*");  
        }  
    }  
}
```

Output: **ow\*is\*the\*time\***





## Additional methods

- If `m` is a matcher, then
  - `m.replaceFirst(replacement)` returns a new String where the first substring matched by the pattern has been replaced by *replacement*
  - `m.replaceAll(replacement)` returns a new String where every substring matched by the pattern has been replaced by *replacement*
  - `m.find(startIndex)` looks for the next pattern match, starting at the specified index
  - `m.reset()` resets this matcher
  - `m.reset(newText)` resets this matcher and gives it new text to examine (which may be a *String*, *StringBuffer*, or *CharBuffer*)





# Some simple patterns

`abc`

exactly this sequence of three letters

`[abc]`

any *one* of the letters `a`, `b`, or `c`

`[^abc]`

any character *except* one of the letters `a`, `b`, or `c`  
(immediately within an open bracket, `^` means “not,”  
but anywhere else it just means the character `^`)

`[a-z]`

any *one* character from `a` through `z`, inclusive

`[a-zA-Z0-9]`

any *one* letter or digit







# Sequences and alternatives

- If one pattern is followed by another, the two patterns must match consecutively
  - For example, `[A-Za-z]+[0-9]` will match one or more letters immediately followed by one digit
- The vertical bar, `|`, is used to separate alternatives
  - For example, the pattern `abc|xyz` will match either `abc` or `xyz`






# Some predefined character classes

- `.` any one character except a line terminator
- `\d` a digit: `[0-9]`
- `\D` a non-digit: `[^0-9]`
- `\s` a whitespace character: `[\t\n\x0B\f\r]`
- `\S` a non-whitespace character: `[^\s]`
- `\w` a word character: `[a-zA-Z_0-9]`
- `\W` a non-word character: `[^\w]`

Notice the space.  
Spaces are **significant**  
in regular expressions!





# Boundary matchers

■ These patterns match the *empty string* if at the specified position:

- ^ the beginning of a line
- \$ the end of a line
- \b a word boundary
- \B not a word boundary
- \A the beginning of the input (can be multiple lines)
- \Z the end of the input except for the final terminator, if any
- \z the end of the input
- \G the end of the previous match





# Greedy quantifiers

(The term “greedy” will be explained later)

Assume  $X$  represents some pattern

$X?$  optional,  $X$  occurs once or not at all

$X^*$   $X$  occurs zero or more times

$X^+$   $X$  occurs one or more times

$X\{n\}$   $X$  occurs exactly  $n$  times

$X\{n, \}$   $X$  occurs  $n$  or more times

$X\{n, m\}$   $X$  occurs at least  $n$  but not more than  $m$  times

Note that these are all *postfix* operators, that is, they come *after* the operand





# Types of quantifiers

- ▶ A greedy quantifier will match as much as it can, and back off if it needs to
  - ▶ We'll do examples in a moment
- ▶ A reluctant quantifier will match as little as possible, then take more if it needs to
  - ▶ You make a quantifier reluctant by appending a `?`:  
`X??`   `X*?`   `X+?`   `X{n}?`   `X{n,}?`   `X{n,m}?`
- ▶ A possessive quantifier will match as much as it can, and never let go
  - ▶ You make a quantifier possessive by appending a `+`:  
`X?+`   `X*+`   `X++`   `X{n}+`   `X{n,}+`   `X{n,m}+`





# Quantifier examples

- ▶ Suppose your text is **aardvark**
  - ▶ Using the pattern **a\*ardvark** (**a\*** is greedy):
    - ▶ The **a\*** will first match **aa**, but then **ardvark** won't match
    - ▶ The **a\*** then “backs off” and matches only a single **a**, allowing the rest of the pattern (**ardvark**) to succeed
  - ▶ Using the pattern **a\*?ardvark** (**a\*?** is reluctant):
    - ▶ The **a\*?** will first match zero characters (the null string), but then **ardvark** won't match
    - ▶ The **a\*?** then extends and matches the first **a**, allowing the rest of the pattern (**ardvark**) to succeed
  - ▶ Using the pattern **a\*+ardvark** (**a\*+** is possessive):
    - ▶ The **a\*+** will match the **aa**, and will not back off, so **ardvark** never matches and the pattern match fails






# Capturing groups

- In regular expressions, parentheses are used for grouping, but they also **capture** (keep for later use) anything matched by that part of the pattern
  - Example: `([a-zA-Z]*)([0-9]*)` matches any number of letters followed by any number of digits
    - If the match succeeds, `\1` holds the matched letters and `\2` holds the matched digits
    - In addition, `\0` holds everything matched by the entire pattern
- Capturing groups are numbered by counting their *opening parentheses* from left to right:
  - `( ( A ) ( B ( C ) ) )`  
1 2 3 4  
`\0 = \1 = ((A)(B(C))), \2 = (A), \3 = (B(C)), \4 = (C)`
- Example: `([a-zA-Z])\1` will match a double letter, such as letter





# Capturing groups in Java

- ▶ If `m` is a matcher that has just performed a successful match, then
  - ▶ `m.group(n)` returns the String matched by capturing group *n*
    - ▶ This could be an empty string
    - ▶ This will be `null` if the pattern as a whole matched but this particular group didn't match anything
  - ▶ `m.group()` returns the String matched by the entire pattern (same as `m.group(0)`)
    - ▶ This could be an empty string
- ▶ If `m` didn't match (or wasn't tried), then these methods will throw an `IllegalStateException`







# Pig Latin

- ▶ Pig Latin is a spoken “secret code” that many English-speaking children learn
  - ▶ There are some minor variations (regional dialects?)
- ▶ The rules for (written) Pig Latin are:
  - ▶ If a word begins with a consonant cluster, move it to the end and add “ay”
  - ▶ If a word begins with a vowel, add “hay” to the end
  - ▶ Example:  
regular expressions are fun! →  
egularray expressionshay arehay unfay!





# Example use of capturing groups

- ▶ Suppose **word** holds a word in English
- ▶ Also suppose we want to move all the consonants at the beginning of **word** (if any) to the end of the word (so **string** becomes **ingstr**)
  - ▶ 

```
Pattern p = Pattern.compile("([aeiou]*)(.*)");  
Matcher m = p.matcher(word);  
if (m.matches()) {  
    System.out.println(m.group(2) + m.group(1));  
}
```
- ▶ Note the use of **(.\*)** to indicate “all the rest of the characters”





# Pig Latin translator

```
■ Pattern wordPlusStuff =  
    Pattern.compile("([a-zA-Z]+)(^[a-zA-Z]*)");  
    Pattern consonantsPlusRest =  
        Pattern.compile("([^aeiouAEIOU]+)([a-zA-Z]*)");  
  
■ public String translate(String text) {  
    Matcher m = wordPlusStuff.matcher(text);  
    String translatedText = "";  
  
    while (m.find()) {  
        translatedText += translateWord(m.group(1)) + m.group(2);  
    }  
    return translatedText;  
}  
  
■ private String translateWord(String word) {  
    Matcher m = consonantsPlusRest.matcher(word);  
    if (m.matches()) {  
        return m.group(2) + m.group(1) + "ay";  
    }  
    else return word + "hay";  
}
```





# Double backslashes

- Backslashes have a special meaning in regular expressions; for example, `\b` means a word boundary
- The Java compiler treats backslashes specially; for example, `\b` in a String or as a char means the backspace character
- Java syntax rules apply first!
  - If you write `"\b[a-z]+\b"` you get a string with backspace characters in it--this is *not* what you want!
  - Remember, you can quote a backslash with another backslash, so `"\\b[a-z]+\\b"` gives the correct string
- Note: if you *read in* a String from somewhere, you are not *compiling* it, so you get whatever characters are actually there





# Additions to the String class

- ▶ All of the following are **public**:
  - ▶ **public** boolean matches(String *regex*)
  - ▶ **public** String replaceFirst(String *regex*,  
String *replacement*)
  - ▶ **public** String replaceAll(String *regex*,  
String *replacement*)
  - ▶ **public** String[] split(String *regex*)
  - ▶ **public** String[] split(String *regex*, int *limit*)
    - ▶ If the limit *n* is greater than zero then the pattern will be applied at most *n* - 1 times, the array's length will be no greater than *n*, and the array's last entry will contain all input beyond the last matched delimiter.
    - ▶ If *n* is non-positive then the pattern will be applied as many times as possible





# Escaping metacharacters

- ▶ A lot of special characters--parentheses, brackets, braces, stars, plus signs, etc.--are used in defining regular expressions; these are called **metacharacters**
- ▶ Suppose you want to search for the character sequence **a\*** (an **a** followed by a star)
  - ▶ **"a\*"**; doesn't work; that means "zero or more **a**s"
  - ▶ **"a\"\***"; doesn't work; since a star doesn't *need* to be escaped (in Java String constants), Java just ignores the **\**
  - ▶ **"a\\\"\***" *does* work; it's the three-character string **a**, **\**, **\***
- ▶ Just to make things even more difficult, it's *illegal* to escape a *non*-metacharacter in a regular expression
  - ▶ Hence, you can't backslash special characters "just in case"





- ▶ There is only one thing to be said about spaces (blanks) in regular expressions, but it's important:

## ▶ *Spaces are significant!*

- ▶ A space stands for a *space*--when you put a space in a pattern, that means to match a space in the text string
- ▶ It's a *really bad idea* to put spaces in a regular expression just to make it look better





# Regular expressions are a language

- ▶ Regular expressions are **not** easy to use at first
  - ▶ It's a bunch of punctuation, not words
  - ▶ The individual pieces are not hard, but it **takes practice** to learn to put them together correctly
  - ▶ Regular expressions form a miniature programming language
    - ▶ It's a different kind of programming language than Java, and requires you to learn new thought patterns
  - ▶ In Java you can't just *use* a regular expression; you have to first create Patterns and Matchers
  - ▶ Java's syntax for String constants doesn't help, either
- ▶ Despite all this, regular expressions bring so **much power and convenience** to String manipulation that they are well worth the effort of learning.
- ▶ After all, many programming tools enable RegEx functionalities








# Thinking in regular expressions

- ▶ The fundamental concept in regular expressions is automatic backtracking
  - ▶ You match the parts of a pattern left to right
    - ▶ Some pattern parts, such as `x` (the letter “x”), `.` (any one character), and `^` (the beginning of the string) are **deterministic**: they either match or don’t match; there are no other alternatives to try
    - ▶ Other pattern parts are **nondeterministic**: they have alternatives, such as `x*` (zero or more letter “x”s), `x+` (one or more letter “x”s), `[aeiou]` (any vowel), and `yes|no` (either “yes” or “no”)
  - ▶ If some part fails to match, you backtrack to the most recent nondeterministic part and look for a different match for that part





# Backtracking examples

- ▶ Search cases for a **[aeiou]s\$**, that is, a vowel followed by an “s” at the end of the string
  - ▶ **[aeiou]** doesn't match **c**
  - ▶ **[aeiou]** matches **a**, **s** matches **s**, **\$** fails
    - ▶ There is no other possible match for **s** in this position
  - ▶ **[aeiou]** doesn't match **s**
  - ▶ **[aeiou]** matches **a**, **s** matches **s**, **\$** succeeds
- ▶ Search Java for **J.\*.+a**
  - ▶ **J** matches **J**, the **.\*** matches **ava**, the **.+** fails
  - ▶ Backtrack to **.\***: The **.\*** matches **av**, the **.+** matches **a**, the **a** fails
  - ▶ Backtrack to **.\***: The **.\*** matches **a**, the **.+** matches **va**, the **a** fails
  - ▶ Backtrack to **.+**: The **.+** matches **v**, the **a** succeeds





# Hazards of regular expressions

- ▶ Regular expressions are complex
  - ▶ They are often used when you cannot guarantee “good” input, so you have to make them fail-safe
- ▶ Backtracking can be extremely expensive
  - ▶ Avoid `.*` and other highly nondeterministic patterns
  - ▶ Test with non-trivial data to make sure your patterns scale
- ▶ Test thoroughly!
  - ▶ Break a complex regular expression into its components, and test each separately
    - ▶ Every pattern is a *program*, and needs to be treated with respect
  - ▶ Pay special attention to edge cases
- ▶ Consider alternatives
  - ▶ Regular expressions are powerful, **but...** If you can get the job done with a few simple String methods, you probably are better off doing it that way

